
Partial Evaluation and Normalisation by Traversals

Work in progress by:

▶ **Daniil Berezun**

Saint Petersburg State University

▶ **Neil D. Jones**

DIKU, University of Copenhagen (prof. emeritus)

INTRODUCTION

The much-studied game semantics for PCF can be thought of as a PCF interpreter.

Ong [?] shows that

a λ -expression M can be evaluated (i.e. normalised) by the algorithm that constructs a traversal of M .

A *traversal* is a sequence of

- ▶ **subexpressions** of M . This is a finite set, whose elements we will call **tokens** (think: M = program, tokens = program points)
- ▶ Any token may have a **back pointers**.

With this approach to normalisation: there is *no need for β -reduction, environments, “thunks” or “closures”* to do the evaluation

START POINT

- ▶ A view of the Oxford normalisation procedure (ONP for short): It is
an interpreter for λ -expressions
- ▶ ONP builds a set of traversal $\mathcal{T}\mathbf{rav}(M)$
 - Let $tr = t_0 \bullet t_1 \bullet \dots \bullet t_n \in \mathcal{T}\mathbf{rav}(M)$ where t_i is a token
 - Syntax-directed inference rules:
based on syntax of the end-token t_n
 - Action: add 0, 1 or more extensions of tr to $\mathcal{T}\mathbf{rav}(M)$. For each,
 - * Add a new token t' , yielding $tr \bullet t'$
 - * Add a back pointer from t'

Data types:

- $tr \in \mathbf{Tr} = \mathbf{Item}^*$
- $\mathbf{Item} = \mathit{subexpression}(M) \times \mathbf{Tr}$

SOME ONP CHARACTERISTICS

Oxford normalization procedure

- ▶ Applies to **simply-typed** λ -expressions
- ▶ Begins by translating M into η -long form
- ▶ Realises the complete head linear reduction of M
- ▶ **Correctness:** by game semantics and categories, using M 's types
- ▶ ONP uses no β -reduction: all is based on subexpressions of M
- ▶ While running, ONP does not use the types of M at all

Goals of this research:

- ▶ Extend **ONP** to **UNP**, for the *untyped* lambda calculus
- ▶ **Partially evaluate** a normaliser with respect to "static" input M .
Use this to **compile** λ -calculus into a *low-level language*.

PARTIAL EVALUATION, BRIEFLY

A partial evaluator is a **program specialiser**. Defining property of *spec*:

$$\forall p \in \text{Programs} . \forall s, d \in \text{Data} . \llbracket \llbracket \text{spec} \rrbracket (p, s) \rrbracket (d) = \llbracket p \rrbracket (s, d)$$

- ▶ Program speedup by **precomputation**.
- ▶ Given program *p* and **static** data *s*, *spec* builds a **residual program**

$$p_s \stackrel{\text{def}}{=} \llbracket \text{spec} \rrbracket p s$$

- ▶ **Staging transformation:**

- $\llbracket p \rrbracket (s, d)$ is a **1 stage** computation
 - $\llbracket \llbracket \text{spec} \rrbracket (p, s) \rrbracket (d)$ is a **2 stage** computation
- ▶ Applications: **compiling**, and **compiler generation**
An old idea: **Semantics directed compiler generation**

WHY PARTIALLY EVALUATE NP

1. The *spec* equation for a normaliser program $\boxed{\text{NP} : \Lambda \rightarrow \text{Traversals}}$

$$\boxed{\forall M \in \Lambda . \llbracket \llbracket \text{spec} \rrbracket (\text{NP}, M) \rrbracket () = \llbracket \text{NP} \rrbracket (M)}$$

2. λ -calculus tradition: M is self-contained.

So **why break normalisation into 2 stages?**

(a) The specialised output $\text{NP}_M = \llbracket \text{spec} \rrbracket (\text{NP}, M)$ can be in a **much simpler language** than λ -calculus.

Our candidate is some **low-level language, LLL**.

(b) 2 stages will be natural for **semantics-directed compiler generation**.

LLL can be an intermediate language to express semantics:

► $\text{NP}_1 = \llbracket \text{spec} \rrbracket \text{NP } M : \Lambda \rightarrow \text{LLL}$

► $\text{NP}_2 = \llbracket - \rrbracket : \text{LLL} \rightarrow \text{Traversals}$ **a semantic function**

HOW TO PARTIALLY EVALUATE ONP

1. **Annotate** parts of normalization procedure as either **static** or **dynamic**.
Variables ranging over
 - (a) **tokens** are **static**, (subexpressions of M ; **finitely many**);
 - (b) **back pointers** are **dynamic**; (**unboundedly many**)
 - (c) so the **traversal** is **dynamic** too.
2. Computations in normalization procedure NP are either **unfolded** or **residualised** (runtime code is generated to **do them later**)
 - ▶ Perform **fully static** computations **at partial evaluation time**.
 - ▶ Operations to build or test a traversal: generate **residual code**.

THE RESIDUAL PROGRAM $\text{ONP}_M = \llbracket \text{spec} \rrbracket \text{ONP } M$

ONP is not quite structurally inductive but it is **semi-compositional**:

Any recursive ONP call has a substructure of M as argument.

Consequences:

- ▶ The partial evaluator can do, at specialisation time,
all of the ONP operations that depend only on M
- ▶ ONP_M performs no operations at all on lambda expressions
- ▶ A specialised program ONP_M contains “residual code”:
 - operations to extend the traversal
 - operations to follow back pointers
- ▶ Subexpressions of M will appear, but are only used as **tokens**:
Tokens are **indivisible**: used as labels and for equality comparisons

STATUS: OUR WORK ON SIMPLY-TYPED λ -calculus

1. We have one version of ONP in `HASKELL` and another in `SCHEME`
2. `HASKELL` version includes: **typing; conversion to eta-long form; the traversal algorithm itself; and construction of the normalised term.**
3. `SCHEME` version: nearly ready to apply automatic partial evaluation.
Plan: use the `UNMIX` partial evaluator (Sergei Romanenko).
4. The LLL output size is only **linearly larger** than M , $|p_M| = O(|M|)$
5. We have also have a handwritten a ***generating extension*** of ONP.

If $p_M = \llbracket \text{ONP-gen} \rrbracket^{scheme}(M)$ then $\forall M . \llbracket M \rrbracket^\Lambda = \llbracket p_M \rrbracket^{LLL}$

Now: LLL = is a tiny subset of `scheme`, so the output p_M is a `SCHEME` program.

MORE WORK FOR SIMPLY-TYPED λ -calculus

Next steps:

- ▶ Produce a generating extension, **automatically**, by **specialising the specialiser to a Λ -traverser**, using UNMIX
- ▶ Redefine LLL formally as a clean stand-alone subset of HASKELL
- ▶ Use HASKELL supercompiler
- ▶ Extend existing approach to **programs with input dynamic data**

STATUS: OUR WORK ON THE UNTYPED λ -calculus

1. *UNP* is a normaliser for $\Lambda^{untyped}$.
2. *UNP* has been done in `HASKELL` and works on a variety of examples.
3. Some traversal items may have **two back pointers**, in comparison: *ONP* uses only one.
4. As *ONP*, *UNP* is also defined **semi-compositionally** by recursion on syntax of λ -expression *M*.
5. By specialising *UNP*, an **arbitrary untyped λ -expression** can be translated to *LLL*.
6. Correctness proof: pending.

TOWARDS SEPARATING PROGRAMS FROM DATA IN Λ

Also we have a one more direction for research:

1. An idea is to regard a **computation of λ -expression M on input d** as a **two-player game between the LLL-codes for M and d** .
2. An interesting example in this case is a usual λ -calculus definition of function *mult* on Church numerals.
3. Amazing fact is that Loops come from out of nowhere:
 - ▶ **Neither `mul` nor the data contain loops**;
 - ▶ But `mul` function is compiled into **an LLL-program with two nested loops**;
 - ▶ We also expect that we can do the computation **entirely without back pointers**.
4. Right now we are trying to express such program-data games in a **communicating** version of *LLL*.

REFERENCES
