

Good afternoon, my name is Daniil and I'm going to present our joint work with professor Neil Deaton Jones work in progress called "Partial Evaluation and Normalization by traversals".

=====
 === slide ??: introduction =====
 =====

The game semantics for PCF can be thought of as a PCF interpreter. In a set of game semantics papers the denotation of an expression is a game strategy. When played, the game results in a traversal.

Ong's recent paper [14] shows that a *simply-typed lambda-expression* M can be *evaluated* (i.e. normalised) by the algorithm that constructs a *traversal* of M .

A *traversal* in this case is justified sequence of subexpressions of M . For brevity, we will call them *tokens*. One can think that M is a *program* and token is a *program point*. Any token may have a *back pointer*, aka justifier, to some other token in history. These pointers are used to lookup some information about λ -expression in history and to find a binamic binder for variables.

Surprisingly, this approach to normalization can be implemented with *none of the traditional implementation techniques* like β -reductions, environments that binds variables to values, closures and thunks for function calls and parameters.

=====
 === slide ??: start point =====
 =====

Now then, it looks very promising to study an *operational* consequences of game semantics. And a start point for our research is an Oxford normalization procedure. We will use an abbreviation *ONP* for name it.

ONP can be thought as a *an interpreter for λ -expressions*. It constructs a set of traversals $\mathcal{T}rav(M)$ for a given λ -expression M each of that is a path in a tree that represents a normal form of λ -expression M .

But what and how?

Consider a traversal tr from $\mathcal{T}rav(M)$ such that is a sequence from t_0 to t_n where each t_i is a *token*.

On each step *ONP* for all such traversals from $\mathcal{T}rav(M)$ adds one or more extensions in $\mathcal{T}rav(M)$ or maybe zero. If zero then it means that traversal tr is already a path in a tree representing an β -normal η -long form of the term.

Each extension is a traversal obtained by concatenation traversal tr with one new token t' that could have a backpointer somewhere in a current view.

Moreover, *ONP* uses inference rules that are entirely based on a syntax of the end-token t_n . In other words, *ONP* is based on *syntax-directed inference rules*.

Thus, we have two main data types: traversal and items. Traversal is just a list of items where an item is a pair: a token and a backpointer.

=====
 === slide ??: SOME ONP CHARACTERISTICS =====
 =====

Now then, let's summarize some important properties of *Oxford normalization procedure*.

First, *ONP* can be applied only to *simply-typed* λ -terms that are in a *η -long form*.

The long form is obtained by η -expanding term fully and replacing the implicit binary application operator of each redex by the *long application operator* @.

Second, *ONP* realises the **complete head linear reduction**. Complete head linear reduction is a repeated applying of *head linear reduction* to all arguments while head linear reduction itself on each step substitutes only one occurrence of variable that is a *hoc-position*.

The correctness of normalization by traversals is done by game semantics and categories, and fully based on *M*'s types.

Then, by contrast to standard approaches to term normalization, method of normalization by traversals uses **no β -reduction** and leaves the original term intact.

Finally, an important property of *ONP* is that while running *ONP* does not use the types of *M* at all. As was mentioned before, *ONP* constructs traversal using only **syntax-directed rules**.

And now now we are able to formulate **goals** of this research:

- First goal is to extend *ONP* to *UNP* that is a normalizer for the *untyped* lambda calculus and
- Second, **Partially evaluate** a normaliser with respect to "static" input λ -term *M*. By partially evaluating normalizer it becomes possible to **compile** λ -calculus into a **low-level language** (LLL for brevity) that further can be used to express **semantics**.

=====
 === slide ?? : PARTIAL EVALUATION, BRIEFLY =====
 =====

Now I want briefly remind what is partial evaluation and how it can be applied to normalizer.

A partial evaluator is a **program specialiser** that can be defined as follows: *spec*:

$\forall p \in \text{Programs} . \forall s, d \in \text{Data} . \llbracket \llbracket \text{spec} \rrbracket(p, s) \rrbracket(d) = \llbracket p \rrbracket(s, d)$

Specializer is a program *spec* that being applied to initial program *p* and data *s* and then to the remaining data *d* is absolutely the same that apply initial program *p* to both data *s* and *d*.

- In other words, this means that for a given program *p* and some static data *s*, partial evaluator builds a **residual program** $p_s \stackrel{\text{def}}{=} \llbracket \llbracket \text{spec} \rrbracket p s \rrbracket$ that being applied to the remaining data *d* produces the same result as an initial program *p* being applied to both of them.
- A net effect of partial evaluation is a **staging program transformation**. I.e. partial evaluation divides computation in two stages: optimized residual program generation and running this program on some dynamic data. While running an initial program *p* on data *s* and *d* is a one stage computation.
- Partial evaluation produces some speed up by **precomputing** all static input at compile time. It can be applied in **compilation**, moreover PE permits an **automatic compiler generation** from an interpreter by self-applying a specializer.

=====
 === slide ??: why partially evaluate ONP =====
 =====

The *spec* equation for a normaliser program **NP** that is just a function from lambda-calculus to Traversals can be seen on the slide.

You can see that there is no external dynamic data and that is a tradition for λ -calculus that λ -term *M* is self-contained.

So a question is **why break normalization into two stages?**

Well, ... , there are several reasons for it:

1. First, a specialiser output $NP_M = \llbracket spec \rrbracket (NP, M)$ can be in a **much simpler language** than λ -calculus. And our candidate for it is some **low-level language** called *LLL*.
2. Second, two stages will be natural for **semantics-directed compiler generation**. Our **aim is to use LLL as an intermediate language to express semantics**. Programs on this low-level language can be thought as a **semantics** for programs from λ -calculus.
 In other words, we **factor** the initial normalization procedure NP into two stages:
 First stage that called NP_1 is a result of partially evaluating normalization procedure to input term *M*
 $NP_1 = \llbracket spec \rrbracket NP\ M : \Lambda \rightarrow LLL$
 and the second stage, NP_2 is the *semantic function* of LLL-programs
 $NP_2 = \llbracket \cdot \rrbracket : LLL \rightarrow Traversals.$

=====
 === slide ??: how to partially evaluate ONP =====
 =====

A next question is "how to partially evaluate oxford normalization procedure with respect to the **static** input term *M*"?

1. First, we have to **annotate** parts of normalization procedure as either **static** or **dynamic**. And here variables ranging over
 - (a) **tokens** that are **static**,
 Because there are only **finitely many** subexpressions of *M*.

 And all other data is actually dynamic
 - (b) i.e. **back pointers** are **dynamic**;
 - (c) and so the **traversal** being built from both of them is **dynamic** too.
2. Computations in normalization procedure **NP** are either **unfolded** by partial evaluator in compile time or **residualised** that means that partial evaluator will generate a runtime code to **do them later** in run-time.
 And then
 - Perform all **fully static** computations **at partial evaluation time**.
 - and for operations to build or test a traversal: generate **residual code**.

==== slide ??: The residual program $ONP_M = \llbracket spec \rrbracket NP\ M$ =====

Now we will talk about structure of a specialized program ONP_M .

Note, that ONP is not quite structually inductive but it is **semi-compositional**:
 in a sence that **Any recursive ONP call has a substructure of M as argument.**
 This property has several **consequences**:

- Firts, the partial evaluator can do, at specialisation time, **all of the ONP operations that depend only on input term M**
- **wherein** this also means that ONP_M performs **no operations at all** on lambda expressions(!)

and for all other operations

- the specialised program ONP_M will be generated. This program contains “residual code”, that means that it contains only operations to build the traversal. There are two kind of operation to do this:
 - operations to extend the traversal; and sometimes
 - operations to follow back pointers when needed to do this
- An important fact is that subexpressions of M will appear, but are only used as **tokens**:
 This means that tokens are **indivisible**: they are only used as labels (i.e. program points) and for equality comparisons with other tokens. Actually we use names instead of real subexpressions. And real subexpressions is only needed for the normalized term reconstruction from traversals.

==== slide ??: Status: our work on simply-typed λ -calculus=====

Status of our work for symply-tiped λ -calculus is as follows:

1. We have one version of ONP written in **Haskell** and another in **Scheme**
2. The Haskell version includes: **typing** using algorithm W ; **conversion to eta-long form; the traversals generation algorithm itself; and construction of the normalised term from the set of traversals.**
3. While Scheme version is nearly ready to apply automatic partial evaluation. We are planning to use the **unmix** partial evaluator that is written by Sergei Romanenko and others. unmix is a general partial evaluator for Scheme. We expect that we will achieve the described above effect of specialising ONP .
4. An important fact is that the **LLL** output program size is only **linearly larger** than M , satisfying

$$|p_M| = O(|M|)$$

while traversal itself can be unboundaly larger than size of the input term.

5. We have also have a handwritten a *generating extension* of *ONP* *ONP – gen*. Symbolically,

$$\text{If } p_M = \llbracket \text{ONP-gen} \rrbracket^{\text{scheme}}(M) \text{ then } \forall M . \llbracket M \rrbracket^\Lambda = \llbracket p_M \rrbracket^{\text{LLL}}$$

It means that by given λ -term *M* *OMP – gen* generates an LLL program p_M that being executed generates a traversal for *M*.

For now: *LLL* is a tiny subset of *scheme*, so the output p_M is a *scheme* program.

6. There are more thing to do for simply-typed λ -calculus:
- First, produce a generating extension *automatically* by *specialising the specialiser to a Λ -traverser* using *unmix*.
 - Second, redefine *LLL* formally as a *clean stand-alone tiny first-order subset* of *haskell* and use *haskell supersompiler* to achive a *partial evaluation* effect.
 - Also we whant to extend existing approach to *programs with input dynamic data* in run-time.

=====
 === Status: our work on the *untyped* λ -calculus ===
 =====
 For *untyped* λ -calculus

1. We have a normaliser *UNP* that is a normaliser for arbitrary untyped lambda term.
2. *UNP* has been done in Haskell and works on a variety of examples. Right now we are working on a more abstract definition of *UNP*.
3. Some of traversal items may have *two back pointers*, in comparison: *ONP* uses only one.
4. As *ONP*, *UNP* is also defined *semi-compositionally* by recursion on syntax of λ -expression *M*. This actually means that it also can be specialized as *ONP* can.
5. Moreover, by specialising *UNP*, an *arbitrary untyped λ -expression* can be translated to our low-level language. So the specialised version of *UNP* could be a semantic function for λ -calculus.
6. Correctness proof: pending. For now, we are working on a correctness proof for *UNP*. We expect that we will prove its correctness formally using some proof assistant like **Coq**.

=====
 === Status: Towards separating programs from data in Λ ===
 Also we have a one more direction for research:

1. An idea is to regard a *computation of λ -expression *M* on input *d** as a *two-player game between the lll-codes for *M* and *d**.

2. An interesting example in this case is a usual λ -calculus definition of multiply function (*mult*) on Church numerals.
3. Amazing fact is that **Loops from out of nowhere**:
 - Neither *mult* nor the data contain loops;
 - but *mult* function is compiled into an *LLL*-program with two nested loops, one to each input numeral, that being applied to two Church numerals computes their product.
 - We expect that we can do the computation entirely without back pointers.
4. Right now we are trying to express such program-data games in a *communicating* version of *LLL*.